



Tertiary Infotech Academy Pte Ltd

UEN: 201200696W

LEARNER GUIDE

For

SQL Fundamental for Beginners

TGS Ref No: TGS-2020505790

Conducted by

Tertiary Infotech Academy Pte Ltd

UEN: 201200696W

Version 13

DOCUMENT VERSION CONTROL RECORD

Version Number	Effective Date of Release	Summary of Included Changes	Author
11.0	16 Oct 2020	Legacy learner materials accompanying the v11 slide deck.	Dr. Alfred Ang
12.0	8 August 2026	Full rebuild in the single-source pipeline: detailed step-by-step for all 11 labs (the slide deck is visual-only), SQL quick reference, glossary, PP + OQ assessment alignment.	Dr. Alfred Ang
13	11 August 2026	Added the SG Mart realistic mock dataset (16 tables, 1,111 rows) with per-lab CSV/Excel/SQL seed folders; rewrote all 11 labs to query business data instead of the generic world database; every SQL statement and every expected result verified against the shipped database; fixed the UNION ORDER BY defect in the full-outer-join lab.	Dr. Alfred Ang

TABLE OF CONTENTS

Introduction.....	4
Course Learning Outcomes.....	4
How You Will Be Assessed.....	4
Before You Start — Environment Setup.....	4
The Course Data — SG Mart Pte Ltd.....	5
Topic 01 — Data Modeling.....	7
Lab 1 — Set Up SQLite Studio & Import the SG Mart Database.....	7
Lab 2 — Create a Database and Tables.....	9
Lab 3 — Model Data with Constraints and Keys.....	13
Topic 02 — Data Processing and Analysis.....	17
Lab 4 — Query Data with SELECT.....	17
Lab 5 — Filter with SQL Operators.....	19
Lab 6 — Insert, Update and Delete Records.....	22
Topic 03 — Data Transformation.....	25
Lab 7 — Aggregate Data with COUNT, AVG and SUM.....	25
Lab 8 — Join Data from Multiple Tables.....	28
Lab 9 — Group Data with GROUP BY and HAVING.....	32
Topic 04 — Introduction to Data Warehouse.....	35
Lab 10 — Map an E-R Model to Database Tables (SMRT Case Study).....	35
Lab 11 — Author Stored Procedures for SQL Server.....	40
SQL Quick Reference.....	43
After the Course.....	44
Glossary.....	44

Introduction

This Learner Guide accompanies the WSQ course SQL Fundamental for Beginners (TGS-2020505790), conducted by Tertiary Infotech Academy Pte Ltd. It provides detailed step-by-step instructions for all 11 hands-on labs, organised by the four course topics: Data Modeling, Data Processing and Analysis, Data Transformation, and Introduction to Data Warehouse. Every lab runs in SQLite Studio (or the free sqliteonline.com) against the SG Mart mock database supplied with the course.

Use this guide alongside the course slides and the lab files in the `courseware/labs/` folder of the course repository. The slides present each concept and lab visually; THIS guide carries the full step-by-step SQL for you to type and run. The guide is an approved open-book reference during the final assessment.

Course Learning Outcomes

- LO1: Apply data modeling for business processes.
- LO2: Apply data processing and analysis using SQL.
- LO3: Apply data transformation from multiple data sources.
- LO4: Apply data mapping to data warehouse.

Skills Framework: Data Engineering (ICT-DIT-3005-1.1), ICT Skills Framework.

How You Will Be Assessed

The final assessment is conducted at the end of the training day and consists of two instruments:

- Practical Performance (PP) — hands-on SQL tasks, 70 minutes, open book, assessor-to-candidate ratio 1:10.
- Oral Questioning (OQ) — 20 minutes, one-to-one with the assessor.
- Format: Open Book — the course slides, this Learner Guide and other approved materials may be used.
- A minimum of 75% attendance is required and learners must be assessed as 'Competent' to be eligible for funding.
- You must take the Assessment digital attendance (TRAQOM) before the assessment begins, and sign the Assessment Summary Record after the Oral Questioning.
- An appeal process is available if required — speak to the trainer or the administrator.

Before You Start — Environment Setup

What you need

- SQLite Studio — free, open-source, cross-platform. Download from <https://sqlitestudio.pl> and install for your OS.

- No-install alternative: the free cloud editor at <https://sqliteonline.com> (also offers an MS SQL engine for the stored-procedure lab).
- The course mock data — download the datasets folder from the course materials on <https://lms-tms.tertiaryinfotech.com>. It contains a prebuilt sgmart.db database, an Excel workbook and CSV files for every lab.
- Any laptop OS — Windows, macOS or Linux.

The Course Data — SG Mart Pte Ltd

Every lab in this course works with one continuous business story so you are always querying something that means something. SG Mart Pte Ltd is a fictitious Singapore retail chain with eight outlets. You will model its shops, staff, members, suppliers and products, then analyse a full year of its sales. Topic 4 switches to an SMRT public-transport model for the data-warehouse case study.

Note: All of this data is invented for training purposes. The names, phone numbers, email addresses and badge numbers do not identify any real person or organisation, and the SMRT figures are illustrative rather than operational records.

The tables you will work with

Table	Rows	What it holds
Outlets	8	The 8 SG Mart retail outlets — code, name, planning area, postal sector, opening date and floor area.
Categories	8	8 product categories grouped into Perishables, Packaged and Non-Food.
Suppliers	9	9 suppliers with contact details, country of origin and lead time in days.
Customers	60	60 loyalty members — tier, join date, points, birth year (some NULL) and home district.
Staff	42	42 employees across the 8 outlets — role, salary, hire date and contact details (some emails are NULL on purpose).
Products	25	25 SKUs with cost, retail price, category, supplier and reorder level (some NULL).
Orders	180	180 sales orders across 2025 — outlet, member (NULL for walk-ins), cashier, channel, payment and status.
OrderItems	685	Order lines (the fact table) — quantity, unit price, discount and line total. Makes SUM/AVG/GROUP BY meaningful.
Routes	8	8 SMRT routes (6 MRT lines + 2 bus services) — the parent

		entity of the transport E-R model.
Stations	14	14 MRT stations with line position, interchange flag and opening date.
Timetables	42	42 timetable rows — first/last trip and frequency per station for weekday, Saturday and Sunday/PH.
DisruptedRoutes	6	6 service disruptions in 2025 — type, details, start time, duration and the alternative arrangements offered.

How the tables relate

- Outlets 1—many Orders — every sale is rung up at one branch.
- Customers 1—many Orders — members link to their purchases; walk-in orders have a NULL CustomerID, which is what makes the LEFT JOIN lab meaningful.
- Staff 1—many Orders — the cashier who served the customer; each staff member belongs to one outlet.
- Orders 1—many OrderItems — the order header and its lines. OrderItems is the fact table: one row per product sold, carrying quantity, price, discount and line total.
- Categories 1—many Products and Suppliers 1—many Products — the two ways the range is classified.
- Routes 1—many Stations and 1—many Timetables; Timetables 1—many DisruptedRoutes — the SMRT chain used in Lab 10.

Where to find it

- Each lab has its own folder under courseware/labs/datasets/ holding just the tables that lab needs, as CSV, plus a seed script.
- seed_sqlite.sql creates and fills every table for that lab in one execution — this is the fastest way to start.
- seed_mysql.sql is the same data in MySQL / SQL Server syntax, used for the stored-procedure lab.
- courseware/labs/datasets/_all/ holds the complete set: SG-Mart-Mock-Data.xlsx (one sheet per table), the prebuilt sgmart.db, and every CSV.

A note on empty values

Some columns are deliberately left empty: a few staff have no work email, some members declined to give a birth year, two products have no reorder level, and seventeen orders were walk-in sales with no membership attached. These gaps are intentional. They are what let you practise IS NULL, see how COUNT() skips NULLs, and understand why a LEFT JOIN returns rows an INNER JOIN would silently drop.

Set up SQLite Studio

Install SQLite Studio, then connect a database: Database → Add a database, browse to a .db file (or click the green '+' to create a new one) and Connect to the database from the left-hand Databases panel. Open the query editor with Tools → Open SQL editor and check the active database shown at the top of the editor.

Conventions used in every lab

- SQL statements are shown in code blocks — type or paste them into the SQL editor and press the Execute (▶ / F9) button.
- SQL keywords are shown in UPPERCASE by convention; SQLite accepts any case.
- Statements end with a semicolon; you can run several statements together.
- Labs 1–10 use SQLite; Lab 11 (stored procedures) uses SQL Server syntax — author it on sqliteonline.com's MS SQL engine.
- If a query fails, read the error message at the bottom of the editor — it names the offending token.

Topic 01 — Data Modeling

Data Sources · Data Modeling · SQL Data Types · Create & Manage Databases (TSC mapping: A1, A2, K1, K2)

Key concepts

- Relational databases — A database is one or more tables; each table is rows (records) and columns (fields) with a fixed row type.
- DDL statements — CREATE / ALTER / DROP define databases, tables, indexes and their columns and data types.
- Constraints — NOT NULL, DEFAULT, UNIQUE, CHECK, PRIMARY KEY and FOREIGN KEY enforce rules on the data.
- Data modeling — Primary and foreign keys link tables into entity relationships that mirror the business process.

Lab 1 — Set Up SQLite Studio & Import the SG Mart Database

Objective: LO1 — identify relevant data sources and set up the SQL workbench (A1, A2).

Goal: Install SQLite Studio, connect it to the SG Mart sample database and explore its tables — your data source for the whole course.

What you'll build

A working SQLite Studio with the SG Mart database connected, showing the Outlets, Products and Categories tables. (Tools: SQLite Studio, sgmart.db mock database, LMS course materials.)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in `courseware/labs/datasets/lab-01-*/`, together with `seed_sqlite.sql` — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Outlets	8	The 8 SG Mart retail outlets — code, name, planning area, postal sector, opening date and floor area.
Products	25	25 SKUs with cost, retail price, category, supplier and reorder level (some NULL).
Categories	8	8 product categories grouped into Perishables, Packaged and Non-Food.

SQLite Studio — SQL Editor

Database: sgmart.db

Execute (F9)

Database

Tools

View

Lab 1 — Set Up SQLite Studio & Import the SG Mart Database

QUERY

```
SELECT OutletCode, OutletName, District, FloorAreaSqm
FROM Outlets
ORDER BY FloorAreaSqm DESC;
```

RESULTS

OutletCode	OutletName	District	FloorAreaSqm
OTL02	SG Mart Tampines Hub	Tampines	1680.5
OTL08	SG Mart Serangoon NEX	Serangoon	1540.0
OTL03	SG Mart Jurong East	Jurong East	1420.0
OTL06	SG Mart Punggol Waterway	Punggol	1310.0
OTL01	SG Mart Orchard	Orchard	1250.0

Query executed successfully · 8 rows returned · sgmart database connected

SQL Fundamental for Beginners · TGS-2020505790

Figure 1 — Lab 1: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. Download SQLite Studio for your OS from the official site and install it. (Alternative: use the free cloud version at sqliteonline.com — no install needed.)

<https://sqlitestudio.pl>

Step 2. Download the course mock data from the LMS. You want the datasets folder — it contains `sgmart.db` (a ready-built SQLite database), an Excel workbook and the CSV files.

<https://lms-tms.tertiaryinfotech.com>

Step 3. Open SQLite Studio, click Database → Add a database, browse to the downloaded sgmart.db and click OK.

Step 4. In the left-hand Databases panel, right-click sgmart and choose Connect to the database.

Step 5. Expand the database tree. You should see the SG Mart tables — Outlets, Categories, Suppliers, Products, Customers, Staff, Orders and OrderItems. Double-click Outlets and open the Data tab to preview its 8 rows.

Step 6. Run your first query to confirm the connection works — list the retail outlets by floor area, largest first.

```
SELECT OutletCode, OutletName, District, FloorAreaSqm  
  
FROM Outlets  
  
ORDER BY FloorAreaSqm DESC;
```

Step 7. Open Outlets.csv from the datasets folder in Excel to see the same data as a spreadsheet. This is the raw form the business hands you before it becomes a table.

Test it

The Databases panel shows the connected sgmart database with the Outlets, Products, Customers, Orders and OrderItems tables; the Data tab displays outlet rows; and the ORDER BY query returns 8 outlets led by SG Mart Tampines Hub (1680.5 sqm).

Note: The same steps are in courseware/labs/lab-01-*.md in the course repository.

Lab 2 — Create a Database and Tables

Objective: LO1 — create and manage databases and tables with SQL DDL (A1, A2).

Goal: Create your own database named 'sgmart_practice' and build the Outlets, Categories, Suppliers and Products tables with CREATE TABLE, choosing an appropriate data type for every column.

What you'll build

A new 'sgmart_practice' database containing four tables whose columns, data types and defaults model the SG Mart retail business. (Tools: SQLite Studio, SQL Editor, lab-02 dataset (CSV + seed script).)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in courseware/labs/datasets/lab-02-*/, together with seed_sqlite.sql — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Outlets	8	The 8 SG Mart retail outlets — code, name, planning area, postal sector, opening date and floor area.
Categories	8	8 product categories grouped into Perishables, Packaged and Non-Food.
Suppliers	9	9 suppliers with contact details, country of origin and lead time in days.
Products	25	25 SKUs with cost, retail price, category, supplier and reorder level (some NULL).

SQLite Studio — SQL Editor

Database: sgmart_practice.db

Execute (F9)

Database

Tools

View

Lab 2 — Create a Database and Tables

QUERY

```

DROP TABLE IF EXISTS Outlets;
CREATE TABLE Outlets (
  OutletCode TEXT PRIMARY KEY,
  OutletName TEXT NOT NULL DEFAULT '',
  District TEXT NOT NULL DEFAULT '',
  PostalSector TEXT NOT NULL DEFAULT '',
  OpenedDate TEXT,
  FloorAreaSqm REAL DEFAULT 0.0
);

```

RESULTS

name	type	notnull	dfit_value	pk
OutletCode	TEXT	0	NULL	1
OutletName	TEXT	1	"	0
District	TEXT	1	"	0
PostalSector	TEXT	1	"	0
OpenedDate	TEXT	0	NULL	0
FloorAreaSqm	REAL	0	0.0	0

● Table Outlets created · schema shown from the Structure tab

SQL Fundamental for Beginners · TGS-2020505790

Figure 2 — Lab 2: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. In SQLite Studio click Database → Add a database, then click the green '+' (plus) icon, name the new database file sgmart_practice and save it. Connect to it from the Databases panel.

Step 2. Open the SQL editor: go to Tools → Open SQL editor. Make sure the active database (top of the editor) is sgmart_practice.

Step 3. Create the Outlets table — the retail branches. Note the TEXT primary key and the REAL column for floor area.

```
DROP TABLE IF EXISTS Outlets;

CREATE TABLE Outlets (
    OutletCode    TEXT PRIMARY KEY,
    OutletName    TEXT NOT NULL DEFAULT '',
    District      TEXT NOT NULL DEFAULT '',
    PostalSector  TEXT NOT NULL DEFAULT '',
    OpenedDate    TEXT,
    FloorAreaSqm  REAL DEFAULT 0.0
);
```

Step 4. Create the Categories table — how SG Mart groups what it sells.

```
DROP TABLE IF EXISTS Categories;

CREATE TABLE Categories (
    CategoryCode  TEXT PRIMARY KEY,
    CategoryName  TEXT NOT NULL DEFAULT '',
    CategoryGroup TEXT NOT NULL DEFAULT 'Packaged'
);
```

Step 5. Create the Suppliers table, including a lead time in days used later for reorder planning.

```
DROP TABLE IF EXISTS Suppliers;

CREATE TABLE Suppliers (
    SupplierID    TEXT PRIMARY KEY,
    SupplierName  TEXT NOT NULL DEFAULT '',
    Address       TEXT,
    Phone         TEXT,
    Country       TEXT NOT NULL DEFAULT 'Singapore',
    LeadTimeDays  INTEGER DEFAULT 14
);
```

Step 6. Create the Products table — the item master. Cost and price are REAL, IsActive is a 0/1 flag.

```
DROP TABLE IF EXISTS Products;

CREATE TABLE Products (
    SKU            TEXT PRIMARY KEY,
    ProductName    TEXT NOT NULL DEFAULT '',
    CategoryCode   TEXT NOT NULL DEFAULT '',
    UnitOfMeasure  TEXT NOT NULL DEFAULT 'each',
    UnitCost       REAL DEFAULT NULL,
    UnitPrice      REAL DEFAULT NULL,
    SupplierID     TEXT,
    ReorderLevel   INTEGER DEFAULT NULL,
    IsActive       INTEGER NOT NULL DEFAULT 1
);
```

Step 7. Load the mock data: open seed_sqlite.sql from this lab's dataset folder and execute the whole script. It fills all four tables. (Or use Tools → Import to load each CSV instead.)

Step 8. Right-click the ssmart_practice database and choose Refresh — the four tables appear. Open each table's Structure tab to verify the columns and data types, then check the row counts.

```
SELECT COUNT(*) AS Outlets FROM Outlets;

SELECT COUNT(*) AS Categories FROM Categories;

SELECT COUNT(*) AS Suppliers FROM Suppliers;

SELECT COUNT(*) AS Products FROM Products;
```

Test it

The ssmart_practice database lists Outlets, Categories, Suppliers and Products; each Structure tab shows the declared columns, data types, defaults and primary keys; and the counts return 8, 8, 9 and 25 rows.

Note: The same steps are in courseware/labs/lab-02-*.md in the course repository.

Lab 3 — Model Data with Constraints and Keys

Objective: LO1 — apply data modeling for business processes with constraints, keys and relationships (A1, A2).

Goal: Build a Persons–PersonOrders mini-model: enforce data rules with NOT NULL, UNIQUE, DEFAULT and CHECK, link the tables with a PRIMARY KEY / FOREIGN KEY relationship, and speed up searches with an index.

What you'll build

A two-table data model (Persons ← PersonOrders) with working constraints, a foreign-key relationship and an index. (Tools: SQLite Studio, SQL Editor, lab-03 dataset (Persons + PersonOrders).)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in courseware/labs/datasets/lab-03-*/ , together with seed_sqlite.sql — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Persons	8	Small Persons table for constraint practice (NOT NULL, UNIQUE, CHECK Age >= 18, DEFAULT City).
PersonOrders	8	Child table for the Persons FK demo — some rows point at persons who exist, one is deliberately unmatched.

SQLite Studio — SQL Editor

Database: sgmart_practice.db

Execute (F9)

Database

Tools

View

Lab 3 — Model Data with Constraints and Keys

QUERY

```

CREATE TABLE Persons (
  ID INTEGER NOT NULL UNIQUE,
  LastName TEXT NOT NULL,
  FirstName TEXT,
  Age INTEGER CHECK (Age >= 18),
  City TEXT DEFAULT 'Singapore',
  PRIMARY KEY (ID)
);

INSERT INTO Persons (ID, LastName, FirstName, Age)
VALUES 1, 'Tan', 'Alice', 30

```

RESULTS

ID	LastName	FirstName	Age	City
1	Tan	Alice	30	Singapore

1 row inserted · City filled from DEFAULT · CHECK and UNIQUE constraints active

SQL Fundamental for Beginners · TGS-2020505790

Figure 3 — Lab 3: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. In the sgmart_practice database's SQL editor, create a Persons table that enforces NOT NULL, UNIQUE, DEFAULT and CHECK rules. SG Mart only issues accounts to adults, so Age must be 18 or over.

```
DROP TABLE IF EXISTS Persons;
```

```
CREATE TABLE Persons (
  ID          INTEGER NOT NULL UNIQUE,
  LastName    TEXT NOT NULL,
  FirstName   TEXT,
  Age         INTEGER CHECK (Age >= 18),
  City        TEXT DEFAULT 'Singapore',
  Email       TEXT,
  PRIMARY KEY (ID)
);
```

Step 2. Prove the constraints work: this insert succeeds and fills City with the default value.

```
INSERT INTO Persons (ID, LastName, FirstName, Age)
```

```
VALUES (1, 'Tan', 'Alice', 30);
```

```
SELECT * FROM Persons;
```

Step 3. Now try to break the rules — each of these statements must FAIL (duplicate ID, NULL LastName, under-age CHECK). Run them one at a time and read the error message.

```
INSERT INTO Persons (ID, LastName, Age) VALUES (1, 'Lim', 25);
```

```
INSERT INTO Persons (ID, Age) VALUES (2, 40);
```

```
INSERT INTO Persons (ID, LastName, Age) VALUES (3, 'Lee', 15);
```

Step 4. Load the rest of the mock people from this lab's dataset — run `seed_sqlite.sql`, or paste these rows. Note Divya Kumar and Priya Nair have no email on file (NULL).

```
INSERT INTO Persons (ID, LastName, FirstName, Age, City, Email) VALUES
(2, 'Lim', 'Bryan', 42, 'Singapore', 'bryan.lim@example.com'),
(3, 'Kumar', 'Divya', 27, 'Singapore', NULL),
(4, 'Wong', 'Cheryl', 35, 'Johor Bahru', 'cheryl.wong@example.com'),
(5, 'Ibrahim', 'Hafiz', 24, 'Singapore', 'hafiz.ibrahim@example.com'),
(6, 'Lee', 'Marcus', 51, 'Singapore', 'marcus.lee@example.com'),
(7, 'Nair', 'Priya', 38, 'Singapore', NULL),
(8, 'Goh', 'Wei Ming', 19, 'Singapore', 'weiming.goh@example.com');
```

Step 5. Create the PersonOrders table whose PersonID column is a FOREIGN KEY referencing Persons — the child table pointing at the parent.

```
DROP TABLE IF EXISTS PersonOrders;
```

```
CREATE TABLE PersonOrders (
    OrderID      INTEGER NOT NULL,
    OrderNumber  INTEGER NOT NULL,
    PersonID     INTEGER,
    OrderDate    TEXT,
    Amount       REAL,
    PRIMARY KEY (OrderID),
    FOREIGN KEY (PersonID) REFERENCES Persons(ID)
```

);

Step 6. Load the order rows. The last one has a NULL PersonID — an order taken before the customer registered.

```
INSERT INTO PersonOrders VALUES

(1, 77895, 3, '2025-02-11', 128.40),
(2, 44678, 3, '2025-03-04', 56.90),
(3, 22456, 1, '2025-03-22', 312.75),
(4, 24562, 1, '2025-05-08', 89.20),
(5, 34764, 6, '2025-06-15', 204.60),
(6, 51230, 2, '2025-07-30', 47.85),
(7, 66120, 7, '2025-08-19', 156.30),
(8, 71984, NULL, '2025-09-02', 63.10);
```

Step 7. Check the relationship holds — list each order beside the person who placed it.

```
SELECT o.OrderNumber, p.FirstName, p.LastName, o.Amount
FROM PersonOrders o
INNER JOIN Persons p ON o.PersonID = p.ID
ORDER BY o.Amount DESC;
```

Step 8. Add a column to an existing table with ALTER TABLE.

```
ALTER TABLE Persons ADD MemberTier TEXT DEFAULT 'Basic';
```

Step 9. Create an index so searches on LastName are fast.

```
CREATE INDEX idx_lastname ON Persons (LastName);
```

Test it

The valid insert appears in Persons with City = 'Singapore'; the three rule-breaking inserts each raise a constraint error; the join returns 7 matched orders led by 312.75; PersonOrders shows a foreign key to Persons in its DDL; and idx_lastname is listed under Indexes.

Note: The same steps are in courseware/labs/lab-03-*.md in the course repository.

Topic 02 — Data Processing and Analysis

SQL Queries · SQL Operators · Insert, Update & Delete (TSC mapping: A3, A5, K4)

Key concepts

- Queries — SELECT retrieves data into a result-set; DISTINCT removes duplicates; WHERE filters rows.
- Operators — AND/OR/NOT, IN, BETWEEN, LIKE wildcards and IS NULL build precise filter conditions.
- Shaping results — ORDER BY sorts, LIMIT caps the row count, and aliases (AS) rename columns and tables.
- Changing data — INSERT INTO adds records, UPDATE modifies them and DELETE removes them — always with WHERE.

Lab 4 — Query Data with SELECT

Objective: LO2 — apply data processing and analysis using SQL queries (A3, A5).

Goal: Query the SG Mart product catalogue with SELECT: retrieve all columns, chosen columns, distinct values and filtered rows, then combine sorting and limits to answer real merchandising questions.

What you'll build

A set of working queries answering questions about SG Mart's products, categories and outlets. (Tools: SQLite Studio, sgmart database, lab-04 dataset.)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in `courseware/labs/datasets/lab-04-*/`, together with `seed_sqlite.sql` — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Products	25	25 SKUs with cost, retail price, category, supplier and reorder level (some NULL).
Categories	8	8 product categories grouped into Perishables, Packaged and Non-Food.
Outlets	8	The 8 SG Mart retail outlets — code, name, planning area, postal sector, opening date and floor area.

SQLite Studio — SQL Editor

Database: sgmart.db

Execute (F9)

Database

Tools

View

Lab 4 — Query Data with SELECT

QUERY

```

SELECT ProductName, UnitPrice
FROM Products
ORDER BY UnitPrice DESC
LIMIT 10;

```

RESULTS

ProductName	UnitPrice
Pineapple Tarts 300g	11.9
Laundry Detergent 2L	11.5
Vanilla Ice Cream 1L	9.8
Shampoo 400ml	9.2
Kopi-O Ground Coffee 500g	8.9
Frozen Prawn Dumpling 400g	8.4

Query executed successfully · 10 rows returned · sorted by UnitPrice DESC

SQL Fundamental for Beginners · TGS-2020505790

Figure 4 — Lab 4: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. In the sgmart database's SQL editor, retrieve every row and column of Products, then only the columns a price list needs.

```
SELECT * FROM Products;
```

```
SELECT ProductName, UnitPrice FROM Products;
```

Step 2. List each category code just once with SELECT DISTINCT — how many distinct categories does the range span?

```
SELECT DISTINCT CategoryCode FROM Products;
```

Step 3. Filter rows with WHERE — every product supplied by Lion City Beverages (SUP05).

```
SELECT SKU, ProductName, UnitPrice
```

```
FROM Products
```

```
WHERE SupplierID = 'SUP05';
```

Step 4. Answer: what are the 10 most expensive products, dearest first?

```
SELECT ProductName, UnitPrice
```

```
FROM Products
```

```
ORDER BY UnitPrice DESC
```

```
LIMIT 10;
```

Step 5. Answer: which 5 outlets have the largest floor area, and in which district?

```
SELECT OutletName, District, FloorAreaSqm
```

```
FROM Outlets
```

```
ORDER BY FloorAreaSqm DESC
```

```
LIMIT 5;
```

Step 6. Combine a filter with a sort — active products under \$5, cheapest first. This is the query a promotions planner actually runs.

```
SELECT ProductName, UnitPrice, CategoryCode
```

```
FROM Products
```

```
WHERE IsActive = 1
```

```
    AND UnitPrice < 5.00
```

```
ORDER BY UnitPrice;
```

Test it

Each query runs without error; the DISTINCT query returns 8 category codes; the top-10 price query is led by Pineapple Tarts 300g at \$11.90; and the largest outlet is SG Mart Tampines Hub (1680.5 sqm).

Note: The same steps are in courseware/labs/lab-04-*.md in the course repository.

Lab 5 — Filter with SQL Operators

Objective: LO2 — analyse data with operators, patterns and sorting (A3, A5).

Goal: Sharpen your WHERE clauses with AND/OR/NOT, IN, BETWEEN, LIKE wildcards, NULL tests and aliases to slice the SG Mart customer and staff data precisely.

What you'll build

A query toolkit covering every SQL operator, answering pattern- and range-based questions on the SG Mart data. (Tools: SQLite Studio, sgmart database, lab-05 dataset.)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in courseware/labs/datasets/lab-05-*/, together with seed_sqlite.sql — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Customers	60	60 loyalty members — tier, join date, points, birth year (some NULL) and home district.
Products	25	25 SKUs with cost, retail price, category, supplier and reorder level (some NULL).
Staff	42	42 employees across the 8 outlets — role, salary, hire date and contact details (some emails are NULL on purpose).

SQLite Studio — SQL Editor

Database: sgmart.db

Execute (F9)

Database

Tools

View

Lab 5 — Filter with SQL Operators

QUERY

```

SELECT ProductName, UnitPrice
FROM Products
WHERE ProductName LIKE 'Frozen%'
ORDER BY ProductName;

```

RESULTS

ProductName	UnitPrice
Frozen Prawn Dumpling 400g	8.4
Frozen Roti Prata 20s	7.3

● Query executed successfully · LIKE 'Frozen%' matched 2 rows

SQL Fundamental for Beginners · TGS-2020505790

Figure 5 — Lab 5: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. Combine conditions with AND — Store Managers earning more than \$6,000 a month.

```

SELECT FirstName, LastName, Role, MonthlySalary
FROM Staff
WHERE Role = 'Store Manager'

AND MonthlySalary > 6000;

```

Step 2. Use OR and NOT — everyone who is NOT a Cashier but works part-time.

```
SELECT FirstName, LastName, Role, Employment
FROM Staff
WHERE NOT Role = 'Cashier'
      AND Employment = 'Part-Time';
```

Step 3. Match a list of values with IN — members on the two premium tiers.

```
SELECT FirstName, LastName, MemberTier, PointsBalance
FROM Customers
WHERE MemberTier IN ('Gold', 'Platinum')
ORDER BY PointsBalance DESC;
```

Step 4. Select a range with BETWEEN — everyday products priced between \$3 and \$6.

```
SELECT ProductName, UnitPrice
FROM Products
WHERE UnitPrice BETWEEN 3.00 AND 6.00
ORDER BY UnitPrice;
```

Step 5. Find patterns with LIKE — every frozen line (the % wildcard matches any characters).

```
SELECT ProductName, UnitPrice
FROM Products
WHERE ProductName LIKE 'Frozen%'
ORDER BY ProductName;
```

Step 6. LIKE also matches mid-string — anything sold in a bottle, whatever the size.

```
SELECT ProductName, UnitOfMeasure
FROM Products
WHERE UnitOfMeasure LIKE '%bottle%';
```

Step 7. Test for missing values with IS NULL — staff with no work email on file — and rename a column with an alias.

```
SELECT FirstName || ' ' || LastName AS StaffName,
```

```
        Role AS JobTitle,  
        Email  
FROM Staff  
WHERE Email IS NULL;
```

Step 8. The opposite test — members who DID give a birth year, so marketing can run a birthday campaign.

```
SELECT FirstName, LastName, BirthYear  
FROM Customers  
WHERE BirthYear IS NOT NULL  
ORDER BY BirthYear  
LIMIT 10;
```

Test it

The IN query returns 16 Gold and Platinum members ordered by points; BETWEEN returns 11 products from \$3.20 to \$5.95; LIKE 'Frozen%' returns the 2 frozen lines; and IS NULL returns the 5 staff with no email, shown under the alias StaffName.

Note: The same steps are in courseware/labs/lab-05-*.md in the course repository.

Lab 6 — Insert, Update and Delete Records

Objective: LO2 — process data by inserting, updating and deleting records (A3, A5).

Goal: Change data safely: INSERT a new product line, UPDATE prices with a WHERE clause, DELETE a discontinued record, and reload the dataset from the provided seed script.

What you'll build

A modified Products table proving you can add, change and remove records — plus a cleanly reloaded database. (Tools: SQLite Studio, sgmart database, lab-06 seed script.)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in courseware/labs/datasets/lab-06-*/, together with seed_sqlite.sql — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Products	25	25 SKUs with cost, retail price,

		category, supplier and reorder level (some NULL).
Customers	60	60 loyalty members — tier, join date, points, birth year (some NULL) and home district.
Orders	180	180 sales orders across 2025 — outlet, member (NULL for walk-ins), cashier, channel, payment and status.

SQLite Studio — SQL Editor

Database: sgmart.db

Execute (F9)

Database

Tools

View

Lab 6 — Insert, Update and Delete Records

```

QUERY
INSERT INTO Products (SKU, ProductName, CategoryCode,
    UnitOfMeasure, UnitCost, UnitPrice, SupplierID)
VALUES ('SKU1026', 'Pandan Chiffon Cake 450g',
    'CAT03', 'each', 4.50, 8.20, 'SUP04');

SELECT SKU, ProductName, CategoryCode, UnitPrice
FROM Products WHERE CategoryCode = 'CAT03';

```

RESULTS

SKU	ProductName	CategoryCode	UnitPrice
SKU1008	Wholemeal Loaf 400g	CAT03	2.8
SKU1009	Kaya Bun 4s	CAT03	3.4
SKU1026	Pandan Chiffon Cake 450g	CAT03	8.2

1 row inserted · IsActive and ReorderLevel took their DEFAULT values

SQL Fundamental for Beginners · TGS-2020505790

Figure 6 — Lab 6: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. SG Mart is adding a new bakery line. Insert it, supplying just the columns you know — the rest take their defaults.

```

INSERT INTO Products (SKU, ProductName, CategoryCode, UnitOfMeasure,
UnitCost, UnitPrice, SupplierID)
VALUES ('SKU1026', 'Pandan Chiffon Cake 450g', 'CAT03', 'each', 4.50, 8.20,
'SUP04');

```

Step 2. Verify the insert — find your new row.

```

SELECT * FROM Products WHERE SKU = 'SKU1026';

```

Step 3. Insert several rows in one statement — three new members joining at the Punggol launch.

```

INSERT INTO Customers (CustomerID, FirstName, LastName, MemberTier,
JoinDate, PointsBalance, BirthYear, District, Phone, Email)

VALUES

    (1061, 'Wei Jie', 'Tan', 'Basic', '2025-11-02', 0, 1996, 'Punggol',
'9123 4567', 'weijie.tan@example.com'),

    (1062, 'Nur Syahirah', 'Osman', 'Basic', '2025-11-02', 50, 1990,
'Punggol', '8234 5678', NULL),

    (1063, 'Divya', 'Menon', 'Silver', '2025-11-03', 1200, 1988, 'Sengkang',
'9345 6789', 'divya.menon@example.com');

```

Step 4. Update records with a WHERE clause — a 10% price rise on every Beverages line. Always SELECT first to see what you are about to change.

```

SELECT ProductName, UnitPrice FROM Products WHERE CategoryCode = 'CAT04';

```

```

UPDATE Products

SET UnitPrice = ROUND(UnitPrice * 1.10, 2)

WHERE CategoryCode = 'CAT04';

```

```

SELECT ProductName, UnitPrice FROM Products WHERE CategoryCode = 'CAT04';

```

Step 5. Update more than one column at once — promote a member and credit bonus points.

```

UPDATE Customers

SET MemberTier = 'Gold',

    PointsBalance = PointsBalance + 500

WHERE CustomerID = 1063;

```

Step 6. Delete one specific record by its primary key. (Never run DELETE without WHERE — it removes every row.)

```

DELETE FROM Products WHERE SKU = 'SKU1026';

```

Step 7. Delete a set of rows with a condition — clear out the cancelled orders.

```

SELECT COUNT(*) FROM Orders WHERE Status = 'Cancelled';

DELETE FROM Orders WHERE Status = 'Cancelled';

```


Step 8. Restore the database to its original state: open seed_sqlite.sql from this lab's dataset folder and execute it. The script drops and recreates every table, so your practice changes are undone.

Step 9. Confirm the reload with row counts.

```
SELECT COUNT(*) AS Products FROM Products;
```

```
SELECT COUNT(*) AS Customers FROM Customers;
```

```
SELECT COUNT(*) AS Orders FROM Orders;
```

Test it

The SKU1026 row appears after the INSERT and is gone after the DELETE; the three new members are added; Beverages prices rise by 10% (Sparkling Water 1.5L goes from 2.35 to 2.59); and after reloading, the counts return 25 products, 60 customers and 180 orders.

Note: The same steps are in courseware/labs/lab-06-*.md in the course repository.

Topic 03 — Data Transformation

Aggregate Data · Join Data · Group Data (TSC mapping: A4, A6, K5, K6)

Key concepts

- Aggregate functions — COUNT, AVG and SUM collapse many rows into one summary value.
- Joins — INNER, LEFT, RIGHT and FULL OUTER joins combine rows from two tables on a matching key.
- Set operations — UNION stacks the result-sets of two SELECTs with matching columns.
- Grouping — GROUP BY aggregates per group; HAVING filters groups the way WHERE filters rows.

Lab 7 — Aggregate Data with COUNT, AVG and SUM

Objective: LO3 — transform data with aggregate functions (A4, A6).

Goal: Summarise SG Mart's sales with COUNT, AVG, SUM, MIN and MAX: count transactions, average basket values, total revenue, and combine aggregates with DISTINCT and WHERE to answer real management questions.

What you'll build

A management summary of SG Mart trading — order counts, average and total revenue, and filtered aggregate answers. (Tools: SQLite Studio, sgmart database, lab-07 dataset.)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in courseware/labs/datasets/lab-07-*/, together with seed_sqlite.sql — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Orders	180	180 sales orders across 2025 — outlet, member (NULL for walk-ins), cashier, channel, payment and status.
OrderItems	685	Order lines (the fact table) — quantity, unit price, discount and line total. Makes SUM/AVG/GROUP BY meaningful.
Products	25	25 SKUs with cost, retail price, category, supplier and reorder level (some NULL).

SQLite Studio — SQL Editor

Database: sgmart.db

Execute (F9)

Database

Tools

View

Lab 7 — Aggregate Data with COUNT, AVG and SUM

QUERY

```
SELECT COUNT(*)           AS AllOrders,
       COUNT(CustomerID) AS MemberOrders
FROM Orders;
```

RESULTS

AllOrders	MemberOrders
180	163

● Query executed successfully · COUNT() skips the 17 NULL walk-in members

SQL Fundamental for Beginners · TGS-2020505790

Figure 7 — Lab 7: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. Count the rows in Orders — how many transactions did SG Mart record in 2025?

```
SELECT COUNT(*) AS TotalOrders FROM Orders;
```

Step 2. COUNT ignores NULLs, and that is useful. Compare these two — the difference is the walk-in customers with no membership.

```
SELECT COUNT(*)          AS AllOrders,
       COUNT(CustomerID) AS MemberOrders
FROM Orders;
```

Step 3. Average a numeric column — the mean value of a sales line.

```
SELECT ROUND(AVG(LineTotal), 2) AS AvgLineValue
FROM OrderItems;
```

Step 4. Total a numeric column — SG Mart's full-year revenue.

```
SELECT ROUND(SUM(LineTotal), 2) AS TotalRevenue
FROM OrderItems;
```

Step 5. MIN and MAX bracket the range — the cheapest and dearest thing on the shelf.

```
SELECT MIN(UnitPrice) AS Cheapest,
       MAX(UnitPrice) AS Dearest
FROM Products;
```

Step 6. Answer: how many distinct payment methods do customers actually use?

```
SELECT COUNT(DISTINCT PaymentMethod) AS PaymentMethods
FROM Orders;
```

Step 7. Answer: how many orders were completed, as opposed to refunded or cancelled?

```
SELECT COUNT(*) AS CompletedOrders
FROM Orders
WHERE Status = 'Completed';
```

Step 8. Put it together — the average basket for online orders only, rounded to cents.

```
SELECT ROUND(AVG(oi.LineTotal), 2) AS AvgOnlineLine,
       COUNT(*)                    AS OnlineLines
```

```
FROM Orders o

INNER JOIN OrderItems oi ON o.OrderID = oi.OrderID

WHERE o.Channel = 'Online';
```

Test it

COUNT(*) returns 180 orders while COUNT(CustomerID) returns 163 (17 walk-ins are NULL); average line value is 13.48; total revenue is 9231.43; prices range from 2.10 to 11.90; and there are 5 distinct payment methods with 165 completed orders.

Note: The same steps are in courseware/labs/lab-07-*.md in the course repository.

Lab 8 — Join Data from Multiple Tables

Objective: LO3 — transform data from multiple sources with joins and unions (A4, A6).

Goal: Combine tables: INNER JOIN orders to outlets and products, see how LEFT JOIN keeps the walk-in orders that have no member, chain three joins together, and emulate a FULL OUTER JOIN in SQLite using LEFT JOIN + UNION.

What you'll build

Joined result-sets linking orders, outlets, customers and products, plus a full-outer-join of two practice tables showing matched and unmatched rows. (Tools: SQLite Studio, sgsmart database, lab-08 dataset.)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in courseware/labs/datasets/lab-08-*, together with seed_sqlite.sql — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Orders	180	180 sales orders across 2025 — outlet, member (NULL for walk-ins), cashier, channel, payment and status.
OrderItems	685	Order lines (the fact table) — quantity, unit price, discount and line total. Makes SUM/AVG/GROUP BY meaningful.
Customers	60	60 loyalty members — tier, join date, points, birth year (some NULL) and home district.
Outlets	8	The 8 SG Mart retail outlets — code, name, planning area, postal sector, opening date and floor area.
Products	25	25 SKUs with cost, retail price,

		category, supplier and reorder level (some NULL).
left_t	4	Practice table for the FULL OUTER JOIN emulation (ids 1-4).
right_t	4	Practice table for the FULL OUTER JOIN emulation (ids 3-6).

SQLite Studio — SQL Editor

Database: sgmart.db

Execute (F9)

Database

Tools

View

Lab 8 — Join Data from Multiple Tables

QUERY

```

SELECT p.ProductName,
       oi.Quantity,
       oi.LineTotal
FROM OrderItems oi
INNER JOIN Products p
  ON oi.SKU = p.SKU
ORDER BY oi.LineTotal DESC
LIMIT 15;

```

RESULTS

ProductName	Quantity	LineTotal
Pineapple Tarts 300g	12	128.52
Shampoo 400ml	12	110.4
Shampoo 400ml	12	104.88
Vanilla Ice Cream 1L	12	99.96
Frozen Prawn Dumpling 400g	12	95.76
Cheddar Slices 250g	12	74.4

Query executed successfully · 15 rows returned · INNER JOIN on SKU

SQL Fundamental for Beginners · TGS-2020505790

Figure 8 — Lab 8: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. Inner join Orders and Outlets on the outlet code — each order with the branch that took it.

```

SELECT o.OrderID, o.OrderDate, ot.OutletName, ot.District
FROM Orders o
INNER JOIN Outlets ot
  ON o.OutletCode = ot.OutletCode
LIMIT 20;

```

Step 2. Join the fact table to the item master — what was actually sold, line by line.

```

SELECT oi.OrderID,
       p.ProductName,

```

```

        oi.Quantity,
        oi.LineTotal
FROM OrderItems oi
INNER JOIN Products p
    ON oi.SKU = p.SKU
ORDER BY oi.LineTotal DESC
LIMIT 15;

```

Step 3. LEFT JOIN keeps every row of the left table even without a match — every order, with the member's name where there is one. Walk-ins show NULL.

```

SELECT o.OrderID, o.OrderDate,
       c.FirstName, c.LastName
FROM Orders o
LEFT JOIN Customers c
    ON o.CustomerID = c.CustomerID
ORDER BY o.OrderID
LIMIT 25;

```

Step 4. Isolate the unmatched rows — this is the standard 'find the orphans' pattern.

```

SELECT COUNT(*) AS WalkInOrders
FROM Orders o
LEFT JOIN Customers c
    ON o.CustomerID = c.CustomerID
WHERE c.CustomerID IS NULL;

```

Step 5. Chain three joins to answer a real question: which outlet sold which product, and for how much?

```

SELECT ot.OutletName,
       p.ProductName,
       ROUND(SUM(oi.LineTotal), 2) AS Revenue
FROM Orders o
INNER JOIN Outlets ot ON o.OutletCode = ot.OutletCode

```

```

INNER JOIN OrderItems oi ON o.OrderID = oi.OrderID
INNER JOIN Products p ON oi.SKU = p.SKU
GROUP BY ot.OutletName, p.ProductName
ORDER BY Revenue DESC
LIMIT 10;

```

Step 6. Create two small practice tables to demonstrate a full outer join.

```

DROP TABLE IF EXISTS left_t;
DROP TABLE IF EXISTS right_t;
CREATE TABLE left_t ( id INTEGER, description TEXT );
CREATE TABLE right_t ( id INTEGER, description TEXT );
INSERT INTO left_t VALUES (1,'left 01'),(2,'left 02'),(3,'left 03'),
(4,'left 04');
INSERT INTO right_t VALUES (3,'right 03'),(4,'right 04'),(5,'right 05'),
(6,'right 06');

```

Step 7. Older SQLite builds have no FULL OUTER JOIN keyword — emulate it with a LEFT JOIN in each direction combined by UNION. Alias every output column: after a UNION, ORDER BY can only refer to the result-set column names, so 'ORDER BY id' fails unless you name the column id.

```

SELECT l.id AS id,
       l.description AS left_desc,
       r.description AS right_desc
FROM left_t l LEFT JOIN right_t r ON l.id = r.id
UNION
SELECT r.id AS id,
       l.description AS left_desc,
       r.description AS right_desc
FROM right_t r LEFT JOIN left_t l ON l.id = r.id
ORDER BY id;

```

Test it

The inner joins return only matched rows; the LEFT JOIN lists every order with NULL names for walk-ins and the orphan count returns 17; the three-way join is led by SG Mart

outlets selling Pineapple Tarts; and the full-join emulation returns ids 1–6 with NULLs on the sides that have no match.

Note: The same steps are in `courseware/labs/lab-08-*.md` in the course repository.

Lab 9 — Group Data with GROUP BY and HAVING

Objective: LO3 — transform data by grouping and filtering groups (A4, A6).

Goal: Aggregate per group: revenue by outlet and by category with GROUP BY, then keep only the groups that matter with HAVING — and see why HAVING exists where WHERE cannot go.

What you'll build

Grouped sales summaries of the SG Mart data — per-outlet and per-category totals, filtered to the groups that matter. (Tools: SQLite Studio, sgmart database, lab-09 dataset.)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in `courseware/labs/datasets/lab-09-*/`, together with `seed_sqlite.sql` — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Orders	180	180 sales orders across 2025 — outlet, member (NULL for walk-ins), cashier, channel, payment and status.
OrderItems	685	Order lines (the fact table) — quantity, unit price, discount and line total. Makes SUM/AVG/GROUP BY meaningful.
Products	25	25 SKUs with cost, retail price, category, supplier and reorder level (some NULL).
Categories	8	8 product categories grouped into Perishables, Packaged and Non-Food.
Outlets	8	The 8 SG Mart retail outlets — code, name, planning area, postal sector, opening date and floor area.

SQLite Studio — SQL Editor

Database: sgmart.db

Execute (F9)

Database

Tools

View

Lab 9 — Group Data with GROUP BY and HAVING

QUERY

```

SELECT o.OutletCode,
       ROUND(AVG(oi.LineTotal),2) AS AvgLine
FROM Orders o
INNER JOIN OrderItems oi ON o.OrderID = oi.OrderID
GROUP BY o.OutletCode
HAVING AVG(oi.LineTotal) > 13
ORDER BY AvgLine DESC;

```

RESULTS

OutletCode	AvgLine
OTL08	16.09
OTL02	14.36
OTL05	14.26
OTL04	14.23
OTL07	13.43

Query executed successfully · HAVING kept 5 of 8 groups

SQL Fundamental for Beginners · TGS-2020505790

Figure 9 — Lab 9: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. Count the orders taken by each outlet — one row per group.

```

SELECT OutletCode, COUNT(*) AS Orders
FROM Orders
GROUP BY OutletCode
ORDER BY Orders DESC;

```

Step 2. Group across a join — total revenue per outlet, with a readable name and alias.

```

SELECT ot.OutletName,
       ROUND(SUM(oi.LineTotal), 2) AS Revenue
FROM Orders o
INNER JOIN Outlets  ot ON o.OutletCode = ot.OutletCode
INNER JOIN OrderItems oi ON o.OrderID   = oi.OrderID
GROUP BY ot.OutletName
ORDER BY Revenue DESC;

```

Step 3. Group by a different dimension — how customers prefer to pay.

```

SELECT PaymentMethod, COUNT(*) AS Orders
FROM Orders
GROUP BY PaymentMethod
ORDER BY Orders DESC;

```

Step 4. Keep only the groups whose average line exceeds \$13 — WHERE cannot filter aggregates, HAVING can.

```

SELECT o.OutletCode,
       ROUND(AVG(oi.LineTotal), 2) AS AvgLine
FROM Orders o
INNER JOIN OrderItems oi ON o.OrderID = oi.OrderID
GROUP BY o.OutletCode
HAVING AVG(oi.LineTotal) > 13
ORDER BY AvgLine DESC;

```

Step 5. Find your loyal customers — members who placed 4 or more orders.

```

SELECT c.FirstName, c.LastName, c.MemberTier,
       COUNT(*) AS Orders
FROM Orders o
INNER JOIN Customers c ON o.CustomerID = c.CustomerID
GROUP BY c.CustomerID, c.FirstName, c.LastName, c.MemberTier
HAVING COUNT(*) >= 4
ORDER BY Orders DESC;

```

Step 6. Use all three clauses together — WHERE filters rows BEFORE grouping, HAVING filters groups AFTER. Revenue by category for completed orders only, keeping the categories above \$1,000.

```

SELECT cat.CategoryName,
       ROUND(SUM(oi.LineTotal), 2) AS Revenue,
       COUNT(*) AS Lines
FROM Orders o
INNER JOIN OrderItems oi ON o.OrderID = oi.OrderID

```

```

INNER JOIN Products p ON oi.SKU = p.SKU

INNER JOIN Categories cat ON p.CategoryCode = cat.CategoryCode

WHERE o.Status = 'Completed'

GROUP BY cat.CategoryName

HAVING SUM(oi.LineTotal) > 1000

ORDER BY Revenue DESC;

```

Test it

The GROUP BY query returns one row per outlet (8 rows, led by OTL08 with 28 orders); revenue by outlet is led by SG Mart Serangoon NEX at 1496.82; the HAVING > 13 query returns the 5 qualifying outlets; and 18 members have 4 or more orders.

Note: The same steps are in courseware/labs/lab-09-*.md in the course repository.

Topic 04 — Introduction to Data Warehouse

Overview of Data Warehouse · Relational Databases · Stored Procedures (TSC mapping: A7, A8, A9, K3, K7)

Key concepts

- Data warehouse — A consolidated, historical, read-mostly database kept separate from operational systems for analysis.
- OLAP vs OLTP — OLAP serves complex analytical queries; OLTP serves fast day-to-day transactions.
- RDBMS & E-R model — The relational model plus entity-relationship design maps business entities to tables and keys.
- Stored procedures — Prepared SQL saved on the server (SQL Server) for reuse — callable with parameters via EXEC.

Lab 10 — Map an E-R Model to Database Tables (SMRT Case Study)

Objective: LO4 — apply data mapping from an E-R model to a data warehouse schema (A7, A8, A9).

Goal: Take the SMRT public-transport E-R model (Routes, Stations, Timetables, DisruptedRoutes) and map it to physical tables: create each entity with its primary key, wire the foreign-key relationships, load the mock operational data and verify the model with joined queries.

What you'll build

A four-table transport schema in SQLite that faithfully implements the SMRT E-R model with PK/FK mappings, loaded with real-shaped route, timetable and disruption data. (Tools: SQLite Studio, sgmart_practice database, lab-10 dataset (Routes, Stations, Timetables, DisruptedRoutes).)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in courseware/labs/datasets/lab-10-*/, together with seed_sqlite.sql — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Routes	8	8 SMRT routes (6 MRT lines + 2 bus services) — the parent entity of the transport E-R model.
Stations	14	14 MRT stations with line position, interchange flag and opening date.
Timetables	42	42 timetable rows — first/last trip and frequency per station for weekday, Saturday and Sunday/PH.
DisruptedRoutes	6	6 service disruptions in 2025 — type, details, start time, duration and the alternative arrangements offered.

SQLite Studio — SQL Editor

Database: sgmart_practice.db

Execute (F9)

Database

Tools

View

Lab 10 — Map an E-R Model to Database Tables (SMRT Case Study)

```

QUERY
SELECT r.route_name, t.station_code,
       d.disrupt_name, d.duration
FROM DisruptedRoutes d
INNER JOIN Timetables t ON d.timetable_id = t.timetable_id
INNER JOIN Routes r     ON t.route_id     = r.route_id
ORDER BY d.disrupt_date;

```

RESULTS

route_name	station_code	disrupt_name	duration
North South Line	NS17	Signal fault at Bishan	45 min
North South Line	NS01	Track point failure	1 hr 20 min
East West Line	EW01	Traction power trip	35 min
Circle Line	CC15	Flash flood at entrance	2 hr
North East Line	NE17	Faulty train withdrawn	25 min
Thomson-East Coast Line	TE14	Communication loss	50 min

Query executed successfully · 3-table join across the SMRT E-R model

SQL Fundamental for Beginners · TGS-2020505790

Figure 10 — Lab 10: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. Review the E-R model before writing any SQL. Routes (route_id PK) 1—many Stations and 1—many Timetables (timetable_id PK, route_id FK); Timetables 1—many DisruptedRoutes (disrupt_no PK, timetable_id FK). Identify each entity's attributes, its primary key, and which column carries the relationship.

Step 2. Create the Routes parent table — the 6 MRT lines and 2 bus services.

```
DROP TABLE IF EXISTS Routes;
```

```
CREATE TABLE Routes (  
    route_id          CHARACTER(4) NOT NULL PRIMARY KEY,  
    route_type        CHARACTER(15) NOT NULL,  
    route_code        VARCHAR(3) NOT NULL,  
    route_name        VARCHAR(125) NOT NULL,  
    route_direction   CHARACTER(10) NOT NULL,  
    remarks           VARCHAR(255)  
);
```

Step 3. Create the Stations table — each station belongs to one route, so route_id is a FOREIGN KEY.

```
DROP TABLE IF EXISTS Stations;
```

```
CREATE TABLE Stations (  
    station_code      CHARACTER(4) NOT NULL PRIMARY KEY,  
    station_name      VARCHAR(40) NOT NULL,  
    route_id         CHARACTER(4) NOT NULL,  
    line_position     INTEGER NOT NULL,  
    is_interchange    INTEGER NOT NULL DEFAULT 0,  
    opened_date       DATE,  
    FOREIGN KEY (route_id) REFERENCES Routes(route_id)  
);
```

Step 4. Create the Timetables child table — first and last trip per station for each service pattern.

```
DROP TABLE IF EXISTS Timetables;
```

```

CREATE TABLE Timetables (
    timetable_id    INTEGER NOT NULL PRIMARY KEY,
    route_id        CHARACTER(4) NOT NULL,
    station_code    CHARACTER(4) NOT NULL,
    frequency        VARCHAR(20) DEFAULT '',
    days_operation  VARCHAR(30) NOT NULL,
    stop_no         NUMERIC(3) NOT NULL,
    first_trip       CHAR(10) NOT NULL,
    last_trip        CHAR(10) NOT NULL,
    FOREIGN KEY (route_id) REFERENCES Routes(route_id)
);

```

Step 5. Create the DisruptedRoutes table — its timetable_id references Timetables, completing the chain.

```

DROP TABLE IF EXISTS DisruptedRoutes;

CREATE TABLE DisruptedRoutes (
    disrupt_no      INTEGER NOT NULL PRIMARY KEY,
    disrupt_date     DATE NOT NULL,
    disrupt_type     VARCHAR(20) NOT NULL,
    disrupt_name     VARCHAR(50) NOT NULL,
    disrupt_details  VARCHAR(300) NOT NULL,
    start_datetime   CHAR(12) NOT NULL,
    duration         CHAR(20) NOT NULL,
    timetable_id     INTEGER NOT NULL,
    alternatives     VARCHAR(3000) NOT NULL,
    FOREIGN KEY (timetable_id) REFERENCES Timetables(timetable_id)
);

```

Step 6. Load the operational data: open seed_sqlite.sql from this lab's dataset folder and execute it. It fills all four tables — 8 routes, 14 stations, 42 timetable rows and 6 recorded disruptions.

Step 7. Verify the parent table loaded and the model reads correctly.

```
SELECT route_id, route_code, route_name, route_type
FROM Routes
ORDER BY route_id;
```

Step 8. Walk one relationship — every station on the North South Line, in line order.

```
SELECT s.station_code, s.station_name, s.line_position, s.is_interchange
FROM Stations s
INNER JOIN Routes r ON s.route_id = r.route_id
WHERE r.route_code = 'NSL'
ORDER BY s.line_position;
```

Step 9. Verify the mapping end-to-end: join all the way from a disruption back to its route.

```
SELECT r.route_name,
       t.station_code,
       t.days_operation,
       d.disrupt_name,
       d.duration
FROM DisruptedRoutes d
INNER JOIN Timetables t ON d.timetable_id = t.timetable_id
INNER JOIN Routes      r ON t.route_id      = r.route_id
ORDER BY d.disrupt_date;
```

Step 10. Now analyse it like a warehouse would — which route type suffers the most disruptions?

```
SELECT r.route_name,
       COUNT(d.disrupt_no) AS Disruptions
FROM Routes r
LEFT JOIN Timetables      t ON r.route_id      = t.route_id
LEFT JOIN DisruptedRoutes d ON t.timetable_id = d.timetable_id
GROUP BY r.route_name
```

```
ORDER BY Disruptions DESC;
```

Test it

The four tables exist with their PK/FK constraints visible in the DDL; Routes returns 8 rows; the NSL station query returns 4 stations in line order; the three-way join returns the 6 disruptions each linked to a route and station; and the LEFT JOIN summary lists every route including those with zero disruptions.

Note: The same steps are in `courseware/labs/lab-10-*.md` in the course repository.

Lab 11 — Author Stored Procedures for SQL Server

Objective: LO4 — operate data-warehouse platform capabilities with reusable stored procedures (A7, A8, A9).

Goal: Write stored procedures in SQL Server syntax against the SG Mart data — a plain procedure and a parameterised one — and understand where they run (SQL Server / MySQL, not SQLite) and why warehouses use them for repeatable batch work.

What you'll build

Stored-procedure scripts (plain and parameterised) ready to run on SQL Server, plus the EXEC calls that invoke them. (Tools: SQL Editor, SQL Server syntax (sqliteonline.com's MS SQL engine, or any SQL Server), lab-11 dataset (`seed_mysql.sql`).)

Mock data for this lab

Load the data before you start. The tables below are supplied as CSV in `courseware/labs/datasets/lab-11-*/`, together with `seed_sqlite.sql` — open that script in the SQL editor and execute it to create and fill every table in one step.

Table	Rows	What it holds
Customers	60	60 loyalty members — tier, join date, points, birth year (some NULL) and home district.
Orders	180	180 sales orders across 2025 — outlet, member (NULL for walk-ins), cashier, channel, payment and status.
Outlets	8	The 8 SG Mart retail outlets — code, name, planning area, postal sector, opening date and floor area.

SQLite Studio — SQL Editor

Database: MS SQL (sqliteonline.com)

Execute (F9)

Database

Tools

View

Lab 11 — Author Stored Procedures for SQL Server

QUERY

```

CREATE PROCEDURE SelectCustomersByTier
    @Tier nvarchar(10)
AS
SELECT CustomerID, FirstName, LastName,
       MemberTier, PointsBalance
FROM Customers
WHERE MemberTier = @Tier
ORDER BY PointsBalance DESC
GO;

EXEC SelectCustomersByTier @Tier 'Platinum'

```

RESULTS

CustomerID	FirstName	LastName	MemberTier	PointsBalance
1050	Muhammad Danial	Osman	Platinum	28127
1013	Grace	Lee	Platinum	24170
1018	Mei Ling	Chan	Platinum	22741
1027	Cheryl	Ng	Platinum	20954
1037	Wei Jie	Koh	Platinum	14935

Procedure created and executed - 7 rows returned for @Tier = 'Platinum'

SQL Fundamental for Beginners · TGS-2020505790

Figure 11 — Lab 11: the SQL editor after running this lab's key statement.

Step-by-step

Step 1. Stored procedures are not available in SQLite — author these against SQL Server. Free option: open sqliteonline.com and switch the engine to MS SQL. Load this lab's seed_mysql.sql to create the Customers, Orders and Outlets tables with the SG Mart data.

Step 2. A stored procedure is a named, saved SQL script the server keeps for you. Write one that returns all members, then run it to save it.

```
CREATE PROCEDURE SelectAllCustomers
```

```
AS
```

```
SELECT CustomerID, FirstName, LastName, MemberTier, PointsBalance
```

```
FROM Customers
```

```
GO;
```

Step 3. Execute the stored procedure by name — no need to retype the query.

```
EXEC SelectAllCustomers;
```

Step 4. Write a parameterised version that filters members by tier through a @Tier parameter. Parameters are what make a procedure reusable.

```
CREATE PROCEDURE SelectCustomersByTier
```

```

    @Tier nvarchar(10)
AS
SELECT CustomerID, FirstName, LastName, MemberTier, PointsBalance
FROM Customers
WHERE MemberTier = @Tier
ORDER BY PointsBalance DESC
GO;

```

Step 5. Execute it, passing the parameter value — then run it again with a different tier to see the same procedure answer a different question.

```

EXEC SelectCustomersByTier @Tier = 'Platinum';
EXEC SelectCustomersByTier @Tier = 'Gold';

```

Step 6. Write a procedure with two parameters — the pattern a nightly warehouse job would use to extract one outlet's takings for a date range.

```

CREATE PROCEDURE OutletSalesByPeriod
    @Outlet nvarchar(5),
    @FromDate date
AS
SELECT o.OutletCode,
       COUNT(*)      AS Orders,
       MAX(o.OrderDate) AS LatestOrder
FROM Orders o
WHERE o.OutletCode = @Outlet
      AND o.OrderDate >= @FromDate
GROUP BY o.OutletCode
GO;

```

Step 7. Run the two-parameter procedure.

```

EXEC OutletSalesByPeriod @Outlet = 'OTL08', @FromDate = '2025-07-01';

```

Step 8. Discuss: why do data warehouses rely on stored procedures? They centralise the logic, run close to the data (no network round-trips), can be granted to users without exposing the base tables, and give you one place to change a rule.

Test it

All three procedures create without error; EXEC SelectAllCustomers returns the full 60-member list; EXEC SelectCustomersByTier @Tier = 'Platinum' returns only the 7 Platinum members ordered by points; and the two-parameter procedure returns one summary row for OTL08.

Note: The same steps are in courseware/labs/lab-11-*.md in the course repository.

SQL Quick Reference

The statements taught in this course, in one place. All run in SQLite unless marked otherwise.

Define structure (DDL)

```
CREATE DATABASE dbname;           -- create a database (file in
SQLite)

CREATE TABLE t (col TYPE constraint, ...); -- create a table

ALTER TABLE t ADD col TYPE;       -- add a column

CREATE INDEX idx ON t (col);       -- speed up searches

DROP TABLE t; DROP DATABASE dbname; -- delete (careful!)
```

Constraints

```
NOT NULL | UNIQUE | DEFAULT 'value' | CHECK (cond)

PRIMARY KEY (col)

FOREIGN KEY (col) REFERENCES parent(pk)
```

Query (DQL)

```
SELECT col1, col2 FROM t;           -- chosen columns ( * = all )

SELECT DISTINCT col FROM t;         -- unique values

SELECT * FROM t WHERE cond;         -- filter rows

cond: = <> > < >= <= AND OR NOT

      IN (v1,v2) BETWEEN a AND b

      LIKE 'a%' IS NULL EXISTS (subquery)

SELECT col AS alias FROM t;         -- rename in the result
```

```
ORDER BY col ASC|DESC    LIMIT n           -- sort and cap
```

Change data (DML)

```
INSERT INTO t (c1,c2) VALUES (v1,v2);
```

```
UPDATE t SET c1 = v1 WHERE cond;           -- ALWAYS use WHERE
```

```
DELETE FROM t WHERE cond;                 -- ALWAYS use WHERE
```

Transform (aggregate · join · group)

```
SELECT COUNT(c), AVG(c), SUM(c) FROM t WHERE cond;
```

```
SELECT ... FROM t1 INNER JOIN t2 ON t1.k = t2.k;
```

```
SELECT ... FROM t1 LEFT  JOIN t2 ON t1.k = t2.k;
```

```
SELECT ... UNION SELECT ...;              -- same columns both sides
```

```
GROUP BY col HAVING AGG(col) cond         -- filter groups, not rows
```

Stored procedures (SQL Server — not SQLite)

```
CREATE PROCEDURE name @param type AS sql_statement GO;
```

```
EXEC name @param = value;
```

After the Course

- Redo every lab from memory in SQLite Studio — fluency comes from typing the SQL yourself.
- Point SQLite Studio at your own data: import a CSV into a table and analyse it with the Topic 2–3 queries.
- The same SQL works in MySQL, PostgreSQL and SQL Server — load `seed_mysql.sql` on a server engine and rerun the labs.
- Continue with the recommended WSQ courses at <https://www.tertiarycourses.com.sg>.
- Complete the TRAQOM feedback survey and download your e-certificate from <https://lms-tms.tertiaryinfotech.com> after being assessed Competent.

Glossary

- **Database** — An organised collection of data — in SQLite, a single .db file holding tables.
- **Table** — A collection of rows sharing the same columns (fields); the unit SQL queries operate on.
- **Row / Record** — One instance of a table's row type — the smallest unit inserted or deleted.

- **Column / Field** — A named, typed slot that every row of a table has.
- **Result-set** — The table of rows returned by a SELECT query.
- **Constraint** — A rule on a column (NOT NULL, UNIQUE, CHECK, DEFAULT, keys) enforced by the engine.
- **Primary key** — The column(s) that uniquely identify each row — UNIQUE + NOT NULL, one per table.
- **Foreign key** — A column referencing another table's primary key, linking child to parent.
- **Index** — A hidden structure that speeds up searches on the indexed column(s).
- **Join** — Combining rows of two tables on matching key values (INNER/LEFT/RIGHT/FULL).
- **Aggregate function** — COUNT, AVG, SUM, MIN, MAX — collapse many rows into one summary value.
- **GROUP BY / HAVING** — Aggregate per group of equal values / filter those groups.
- **Data modeling** — Designing entities, attributes and key relationships to mirror a business process.
- **E-R model** — Entity-Relationship model — the high-level diagram mapped into tables and keys.
- **Data warehouse** — A separate, consolidated, historical, read-mostly database for analysis and decisions.
- **OLAP / OLTP** — Analytical processing over the warehouse / transactional processing in operational systems.
- **RDBMS** — Relational Database Management System — the engine class SQL was built for.
- **Stored procedure** — Prepared SQL saved on the server (SQL Server/MySQL), executed by name with EXEC.
- **SQL injection** — Attacking a system by smuggling SQL through user input — prevented with parameterised queries.